

Creating a Keras model

INTRODUCTION TO DEEP LEARNING IN PYTHON



Dan Becker

Data Scientist and contributor to Keras
and TensorFlow libraries

Model building steps

- Specify Architecture
- Compile
- Fit
- Predict

Model specification

```
import numpy as np
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential

predictors = np.loadtxt('predictors_data.csv', delimiter=',')
n_cols = predictors.shape[1]


model = Sequential()
model.add(Dense(100, activation='relu', input_shape = (n_cols,)))
model.add(Dense(100, activation='relu'))
model.add(Dense(1))
```

Let's practice!

INTRODUCTION TO DEEP LEARNING IN PYTHON

Compiling and fitting a model

INTRODUCTION TO DEEP LEARNING IN PYTHON



Dan Becker

Data Scientist and contributor to Keras and TensorFlow libraries

Why you need to compile your model

- Specify the optimizer
 - Many options and mathematically complex
 - "Adam" is usually a good choice
- Loss function
 - "mean_squared_error" common for regression

Compiling a model

```
n_cols = predictors.shape[1]
model = Sequential()
model.add(Dense(100, activation='relu', input_shape=(n_cols,)))
model.add(Dense(100, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mean_squared_error')
```

What is fitting a model

- Applying backpropagation and gradient descent with your data to update the weights
- Scaling data before fitting can ease optimization

Fitting a model

```
n_cols = predictors.shape[1]
model = Sequential()
model.add(Dense(100, activation='relu', input_shape=(n_cols,)))
model.add(Dense(100, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mean_squared_error')
model.fit(predictors, target)
```

Let's practice!

INTRODUCTION TO DEEP LEARNING IN PYTHON

Classification models

INTRODUCTION TO DEEP LEARNING IN PYTHON



Dan Becker

Data Scientist and contributor to Keras
and TensorFlow libraries

Classification

- `'categorical_crossentropy'` loss function
- Similar to log loss: Lower is better
- Add `metrics = ['accuracy']` to compile step for easy-to-understand diagnostics
- Output layer has separate node for each possible outcome, and uses `'softmax'` activation

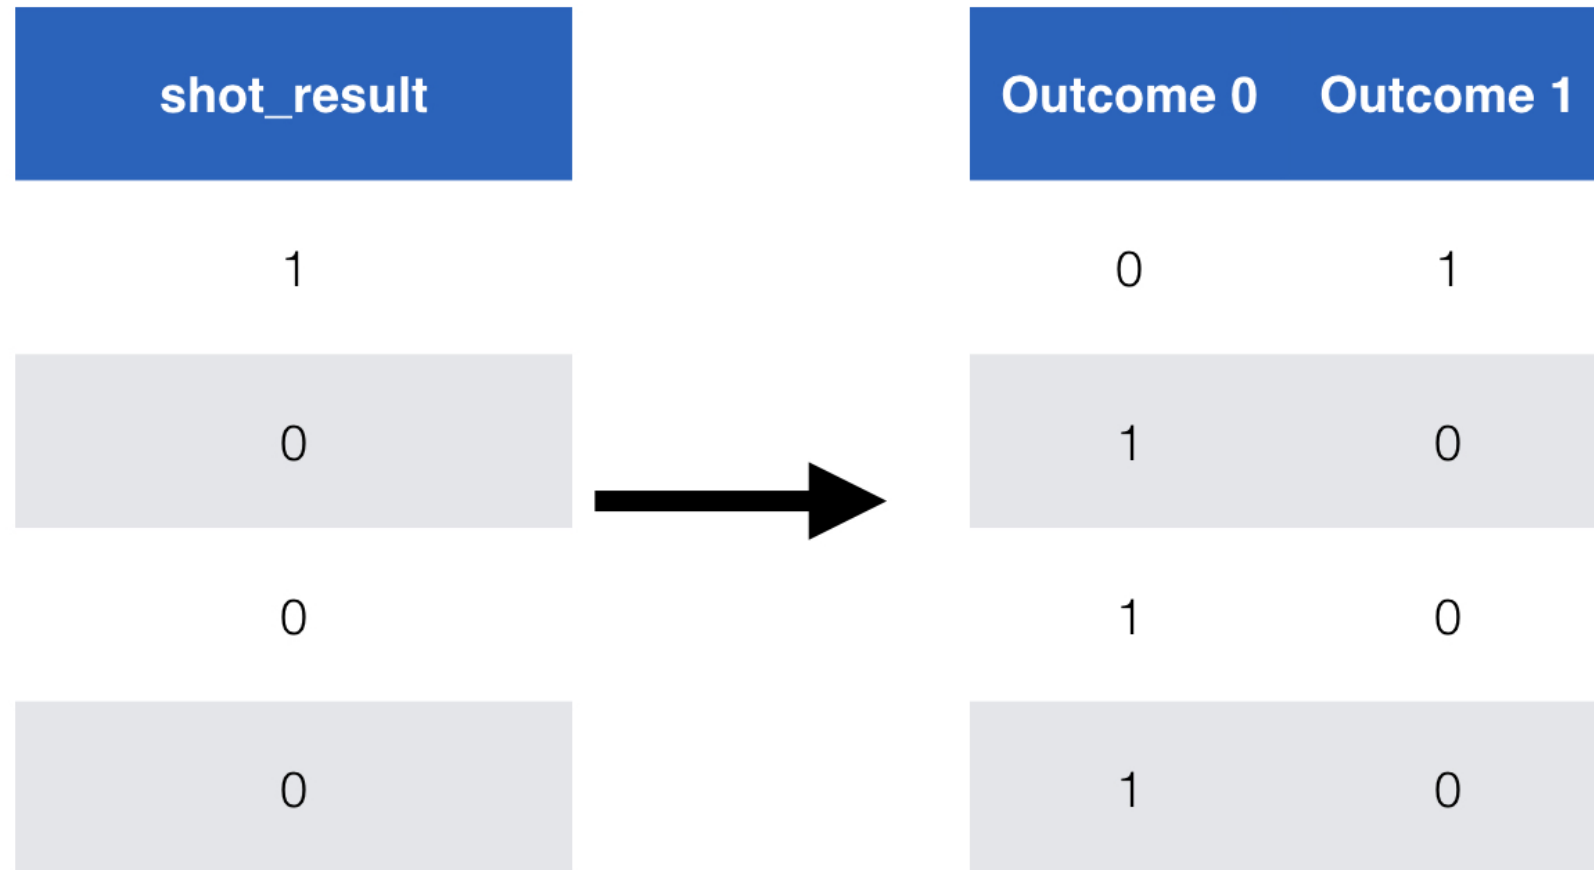
Quick look at the data

shot_clock	dribbles	touch_time	shot_dis	close_def_ dis	shot_result
10.8	2	1.9	7.7	1.3	1
3.4	0	0.8	28.2	6.1	0
0	3	2.7	10.1	0.9	0
10.3	2	1.9	17.2	3.4	0

Quick look at the data

shot_clock	dribbles	touch_time	shot_dis	close_def_dis	shot_result
10.8	2	1.9	7.7	1.3	1
3.4	0	0.8	28.2	6.1	0
0	3	2.7	10.1	0.9	0
10.3	2	1.9	17.2	3.4	0

Transforming to categorical



Classification

```
from tensorflow.keras.utils import to_categorical

data = pd.read_csv('basketball_shot_log.csv')
predictors = data.drop(['shot_result'], axis=1).values
target = to_categorical(data['shot_result'])

model = Sequential()
model.add(Dense(100, activation='relu', input_shape = (n_cols,)))
model.add(Dense(100, activation='relu'))
model.add(Dense(100, activation='relu'))
model.add(Dense(2, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(predictors, target)
```


Classification

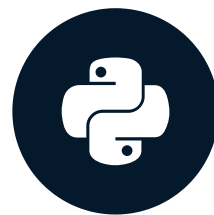
```
Epoch 1/10
128069/128069 [=====] - 4s - loss: 0.7706 - acc: 0.5759
Epoch 2/10
128069/128069 [=====] - 5s - loss: 0.6656 - acc: 0.6003
Epoch 3/10
128069/128069 [=====] - 6s - loss: 0.6611 - acc: 0.6094
Epoch 4/10
128069/128069 [=====] - 7s - loss: 0.6584 - acc: 0.6106
Epoch 5/10
128069/128069 [=====] - 7s - loss: 0.6561 - acc: 0.6150
Epoch 6/10
128069/128069 [=====] - 9s - loss: 0.6553 - acc: 0.6158
Epoch 7/10
128069/128069 [=====] - 9s - loss: 0.6543 - acc: 0.6162
Epoch 8/10
128069/128069 [=====] - 9s - loss: 0.6538 - acc: 0.6158
Epoch 9/10
128069/128069 [=====] - 10s - loss: 0.6535 - acc: 0.6157
Epoch 10/10
128069/128069 [=====] - 10s - loss: 0.6531 - acc: 0.6166
```

Let's practice!

INTRODUCTION TO DEEP LEARNING IN PYTHON

Using models

INTRODUCTION TO DEEP LEARNING IN PYTHON



Dan Becker

Data Scientist and contributor to Keras
and TensorFlow libraries

Using models

- Save
- Reload
- Make predictions

Saving, reloading, and using your Model

```
from tensorflow.keras.models import load_model
model.save('model_file.h5')
my_model = load_model('model_file.h5')
predictions = my_model.predict(data_to_predict_with)
probability_true = predictions[:,1]
```

Verifying model structure

```
my_model.summary()
```

```
-----
Layer (type)                 Output Shape              Param #   Connected to
-----
dense_1 (Dense)              (None, 100)               1100      dense_input_1[0][0]
-----
dense_2 (Dense)              (None, 100)               10100     dense_1[0][0]
-----
dense_3 (Dense)              (None, 100)               10100     dense_2[0][0]
-----
dense_4 (Dense)              (None, 2)                  202       dense_3[0][0]
=====
Total params: 21,502
Trainable params: 21,502
Non-trainable params: 0
```

Let's practice!

INTRODUCTION TO DEEP LEARNING IN PYTHON